

Simulated Annealing

Winter 2024

Population versus Individual



Hill Climb

- The simplest individual-based search method
- Start at a random solution; consider various moves within the neighborhood; accept the ones when advantageous; repeat until no more move can be made



Pseudocode of Hill Climbing

Select a starting solution $s_0 \in S$

While termination condition is not met:

iteration limit or $f(s_0) < f(s)$ for all $s \in N(s_0)$

Select s such that $f(s) < f(s_0)$


Replace s_0 by s

Return s_0 as the approximation to the optimal solution

*** For minimization problem*


Disadvantages of hill climbing

- It is very likely to get stuck in local optimal solutions
- Solution quality is highly sensitive to the initial solution (P always leads to Q; R can be improved a lot more to ★)
- Hard to decide where to go if all the neighbors have the same objective value (After T, where to go?)



Possible Improvements

- Hill climbing in inner-loop and re-start the initialization in outer-loop
- Increase the range (vision) of neighborhood
- Allow “bad move” whose frequency is governed by a probability function



Simulated Annealing: Thermodynamics

- Metropolis's work in statistical thermodynamics
- At temperature t , the probability of an increase in energy of magnitude ΔE is given by

$$P(\Delta E) = \exp(-\Delta E / \alpha t)$$

where α is Boltzmann's constant

1.38065×10^{-23} joules/kelvin



Dr. Nicolas Metropolis
1915-1999




Dr. Ludwig Boltzmann
1844-1902

Simulated Annealing



- If solid material is heated past its melting point and then cooled back into a solid state, the structural property depends on the cooling rate
- The heat causes the **atoms** to become unstuck from their initial positions (a local minimum of the internal energy) and **wander randomly** through states of higher energy
- Slow cooling gives them more chances of finding configurations with lower internal energy than the initial one



Physical Analogy

- If the cooling is done too fast, the resulting solid will be frozen into a locally optimal structure, such as a glass or a crystal with several defects in the structure.
- Inspiration: Finding low energy states of a solid by initially melting the substance, and then lowering the temperature slowly



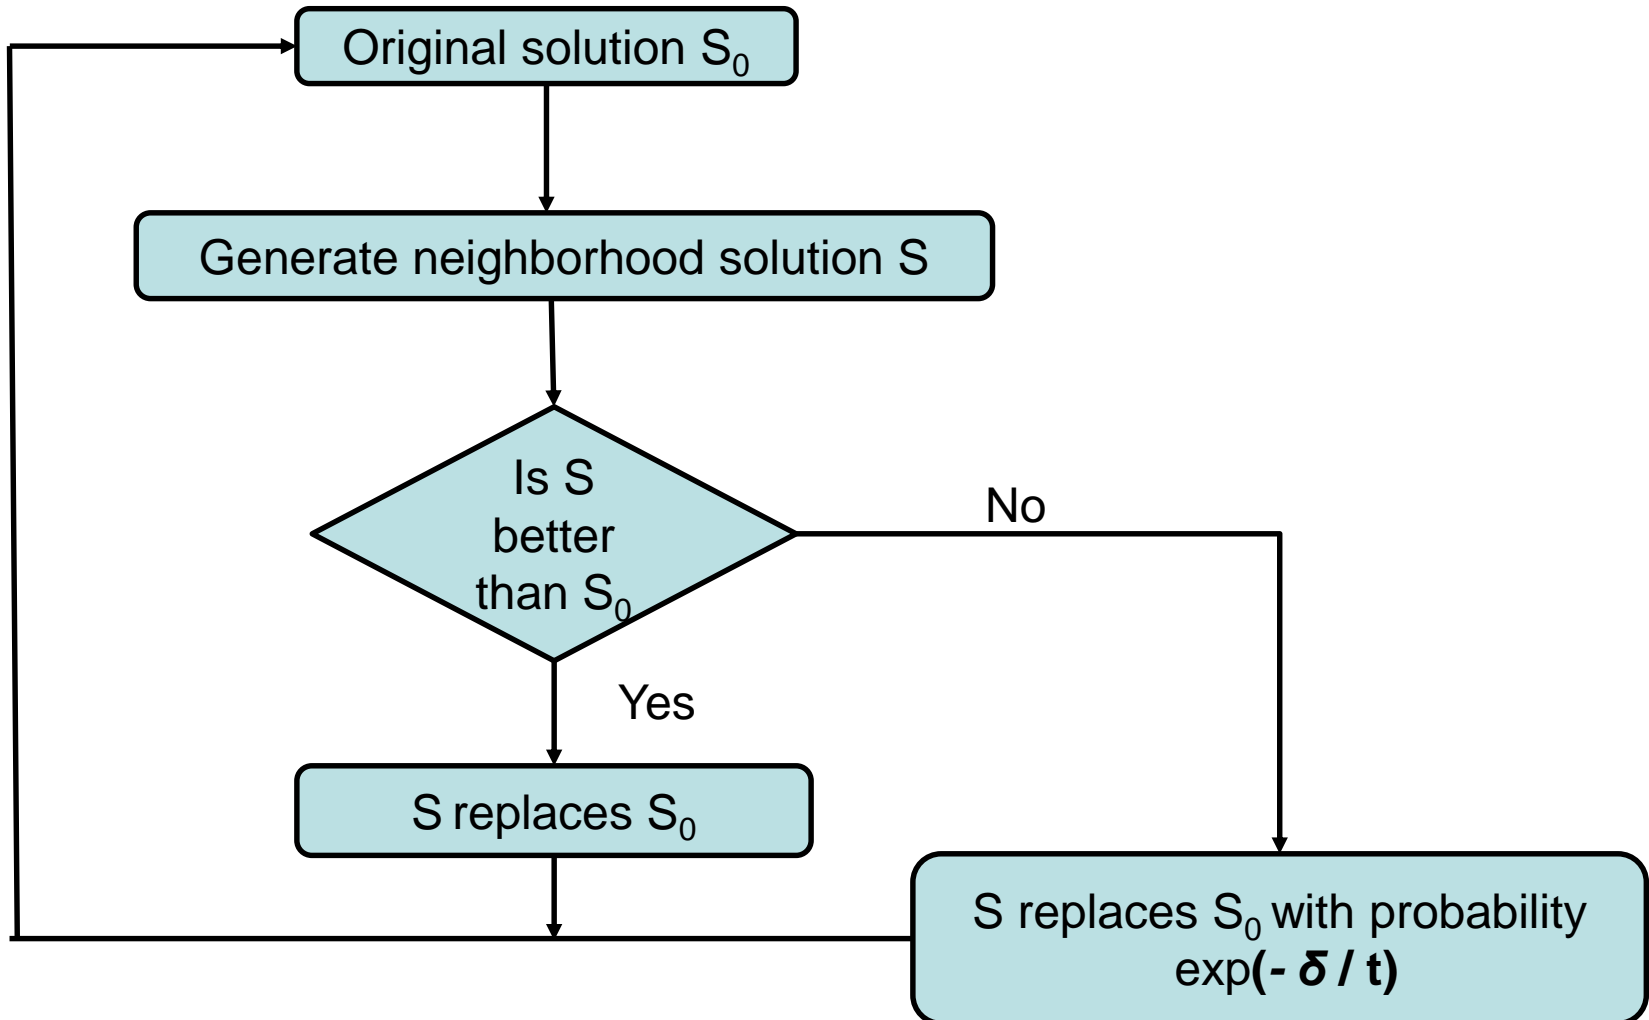
Concept



Prof. Kirkpatrick
Computer Science

- Kirkpatrick, Gelatt, Jr., and Vecchi, (1983) "Optimization by Simulated Annealing", *Science*, 220(4598), 671-680.
- Each step of the Simulated Annealing (SA) algorithm replaces the current solution by a random "nearby" solution, chosen with a probability that depends on **the difference between the corresponding function values** and **on a global parameter T (temperature)**, that is **gradually decreased during the process**.

Flowchart



Pseudocode of SA (1)

Select an initial solution s_0 ; initial temperature t_0

Select a temperature reduction function (cooling schedule) $T(k)$: k is the reduction step

Outer: While temperature $>$ final_temperature

Inner: While iteration limit n_{iter} is not reached

Select $s \in N(s_0)$

$$\delta = f(s) - f(s_0)$$

If $\delta < 0$ (*for minimization problem*) $\delta > 0$ for maximization

$s_0 = s$ **#always accept better moves**

Pseudocode of SA (2)

Else **#may accept worse moves**

generate $r \sim U(0,1)$

If $r < \mathbf{\exp(-\delta / t)}$, $s_0 = s$ (*for minimization*)

If $r < \mathbf{\exp(\delta / t)}$, $s_0 = s$ (*for maximization*)

End **Inner**


$t = T(k)$

End **Outer**


Return s_0

Probability of Adverse Move

- Probability $\exp(-\delta / t)$ is related to δ and t
 - larger $\delta \rightarrow$ smaller acceptance probability
(less likely to accept worse move)
 - lower temperature \rightarrow smaller acceptance probability
(decreasing probability as temperature drops)



$t = 500$



$\delta = 100$

Key issues

For worse moves, if $r < \exp(-\delta / t)$, $s_0 = s$

- Decreasing temperature \rightarrow Decrease the probability of accepting bad moves
- Allow a bad move but the probability of accepting it is **inverse proportional** to the difference from the previous move
- Increasing the number of iterations at each temperature \rightarrow spend **more efforts in later iterations**

Numerical Example

Maximize $f(x)=x^3-60x^2+900x+100$; $x=[0\sim31]$

Starting temperature $t_0=500$


Cooling schedule $T(k)=t_0 0.9^k$; $n_{\text{iter}}=1$ (1 iteration at each temperature)

x is expressed as 5-bit representation; its neighborhood
is selected based on one-bit change

Initial random solution: (10011)

Initial x : $2^4+2^1+2^0=19$

Initial $f(x)$: 2399



Experiment Result

If $r < \exp(\delta / t)$, $s_0 = s$

$\exp(-112/500)=0.79932$

T(k)	Selected bit	String	x	f(x)	δ	Move Prob.	r <i>random</i>	Move?
500	1	00011	3	2287	-112	0.79932	0.638	YES
450	3	00111	7	3803	>0			
405	5	00110	6	3556	-247	0.54342	0.124	YES
364.5	2	01110	14	3684	>0			
328.0	4	01100	12	3988	>0			
295.2	3	01000	8	3972	-16	0.94724	0.837	YES
265.7	4	01010	10	4100	>0			
239.1	5	01011	11	4071	-29	0.88578	0.419	YES
215.2	1	11011	27	343	-3728	3.00E-08	0.573	NO
193.7	2	00011	3	2287	-1784	0.0001	0.249	NO

$\exp(-247/405)=0.54342$

Optimal

Move away

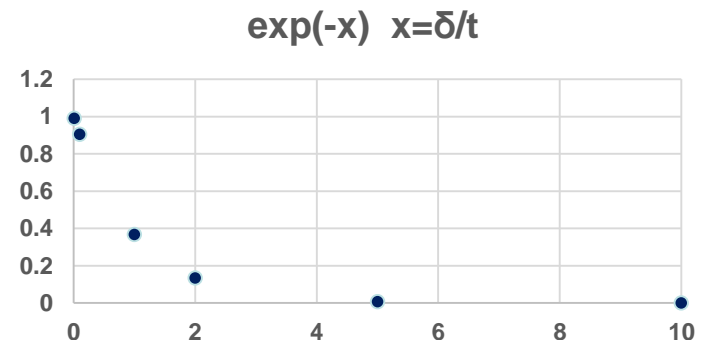
Generic Decisions

- Initial temperature
- Final temperature
- Cooling schedule
 - **Decreasing** function
- Number of iteration at each temperature: n_{iter}
 - Spend longer time at lower temperatures
 - **Increasing** function

Initial Temperature

For worse moves, if $r < \exp(-\delta / t)$, $s_0 = s$

- If **too hot**, we are conducting a pure random search (always accept worse moves)
- If **too cold**, we are implementing hill climbing (do not allow worse moves at all)
- Preliminary run: Start high, freeze quickly until about 50%~80% of worse moves are accepted; use this as the initial temperature



Final Temperature

- It is usual to let the temperature decrease until it is close to zero
- Sufficient to stop when the chance of accepting a bad move is negligible

Cooling Schedule (1)

- Geometrical

$T(k) = t_0 a^k$; a is a constant smaller than (close to) 1

- Logarithmic

$$T(k) = t_0 \ln(2)/\ln(k) \quad k \geq 2$$

- Exponential

$T(k) = t_0 \exp(-c/k)$; c is a small positive constant close to 0

Cooling Schedule (2)

- Lundy and Mees (1986)

$T(k) = T(k-1) / (1 + \gamma T(k-1))$; γ is a small constant

- Aarts and Korst (1989)

$T(k) = T(k-1) / [1 + (T(k-1) \ln(1 + \gamma) / 3\sigma_t)]$;

σ_t is the standard deviation at temperature t ;

γ is a constant

Number of Iterations (1)

- Allow enough iterations at each temperature so that the system stabilises at that temperature
- The number of iterations may vary at different temperatures

Number of Iterations (2)

- A fixed constant
- Linearly increasing
 $n_{\text{iter}} = \alpha(t_0 - t) + \beta$; α and β are constants
- Other **increasing functions**
- When a certain number of better solutions have been found, stop the iterations; go to next temperature

Number of Iterations (3)

- Compromise:
 - either doing a **large** number of iterations at **a few** temperatures
 - or a **small** number of iterations at **many** temperatures
 - a balance between the two

Variant:

Acceptance Probability

- Simplification on the distribution

1. $P(t) = 1 - \delta / t ; \delta < t$

2. Discrete options

Iterations 1~100: 0.8





Iterations 101~200: 0.6

Iterations 201~300: 0.4

Variant:


Alternative Cooling Schedules

- T_i is the temperature for step i , where i increases to N
- The initial and final temperatures, T_0 and T_N



Variant: Neighborhood (1)

- Neighborhood definition is implicit, not shown in the algorithm
- Real-valued neighborhood: analogous to **real-valued mutation**
- Discrete neighborhood:
 - probability of choosing the next solution may be in inverse relation to the distance from the current solution



Variant: Neighborhood (2)

- Permutation neighborhood:

- Inversion

[1 3 4 2 6 5]

[1 4 3 2 6 5]

- Transposition

[1 3 4 2 6 5]

[1 6 4 2 3 5]

- Displacement

[1 3 4 2 6 5]

[1 3 6 4 2 5]

- 2-opt (Croes, 1958)

[1 3 4 2 6 5]

[1 6 2 4 3 5]

Variant: Fixed Temperature

- Simulated annealing does most of its work during cooling stage
- Annealing at a constant temperature
 - Searching for the optimum temperature, which determines the acceptance probability
 - Use the acceptance probability throughout the process

Variant: Reheating

- When a move is accepted, cool the system by $t = t / (1+\beta)$; β is a constant >0
- When a move is rejected, heat the system by $t = t / (1-\gamma)$; γ is a constant >0

following example, $\beta=0.2$ and $\gamma=0.1$, original string=[10011]

T(k)	Selected bit	String	x	f	δ	Move Prob.	Move?	String
500	1	00011	3	2287	-112	0.799315	YES	00011
416.7	3	00111	7	3803	>0		YES	00111
347.2	5	00110	6	3556	-247	0.543418	NO	00111
385.8	2	01111	15	3475	-328	0.427340	YES	01111

$416.7/(1+0.2)$


$347.2/(1-0.1)$

Variant: Objective Function

- The objective function is calculated at every iteration of the algorithm
- If calculation is expensive (e.g., simulation, finite element method)
 - Approximate the objective value (e.g., using fewer simulation runs); focus only on potentially good solutions
 - Evaluate just the changes and update the previous objective value
 - Keep cache to store objective values that have already been evaluated; may need a large memory space

Parallel Implementation

- Independent runs on **different computing processors**, with limited communication
- **Split the search space into parts**, each to be assigned to a particular processor
- Each processor checks the solution and test for acceptance independently. An accepted solution is conveyed to all the other processors



Hybrid with Other Algorithms

- Start with a solution that has been heuristically built by greedy algorithm or other meta-heuristics, such as GA or PSO
- Switch between different algorithms iteratively; usually use GA or PSO for global search and SA for local search
- Use SA to fine tune the control parameters of other algorithms

GA: crossover rate, mutation rate

PSO: w , c_1 , c_2 , and V_{\max}

Summary of SA

- Individual-based, not population-based
- Simple in coding
- Use a nested loop structure
- No memory of past experience
- More complicated to fine tune the algorithm; particularly decreasing cooling schedule and increasing function of iterations